A Random Matrix Perspective of Echo State Networks: From Precise Bias–Variance Characterization to Optimal Regularization

Yessin Moahker[1], Malik Tiomoko[1], Cosme Louart[2], Zhenyu Liao[3]

[1]Huawei Noah's Ark Lab [2]Chinese University of Hong Kong, Shenzhen [3]HUST, Wuhan

1. Introduction: The Echo State Paradigm

Standard RNNs face a **computational nightmare**: vanishing gradients and high energy costs during backpropagation.

Echo State Networks (ESNs) solve this by splitting the architecture:

- **Reservoir (Fixed)**: A high-dimensional nonlinear dynamical system that projects input into a rich feature space without training.
- **Readout (Trained)**: A linear layer trained to **simple Ridge Regression**.

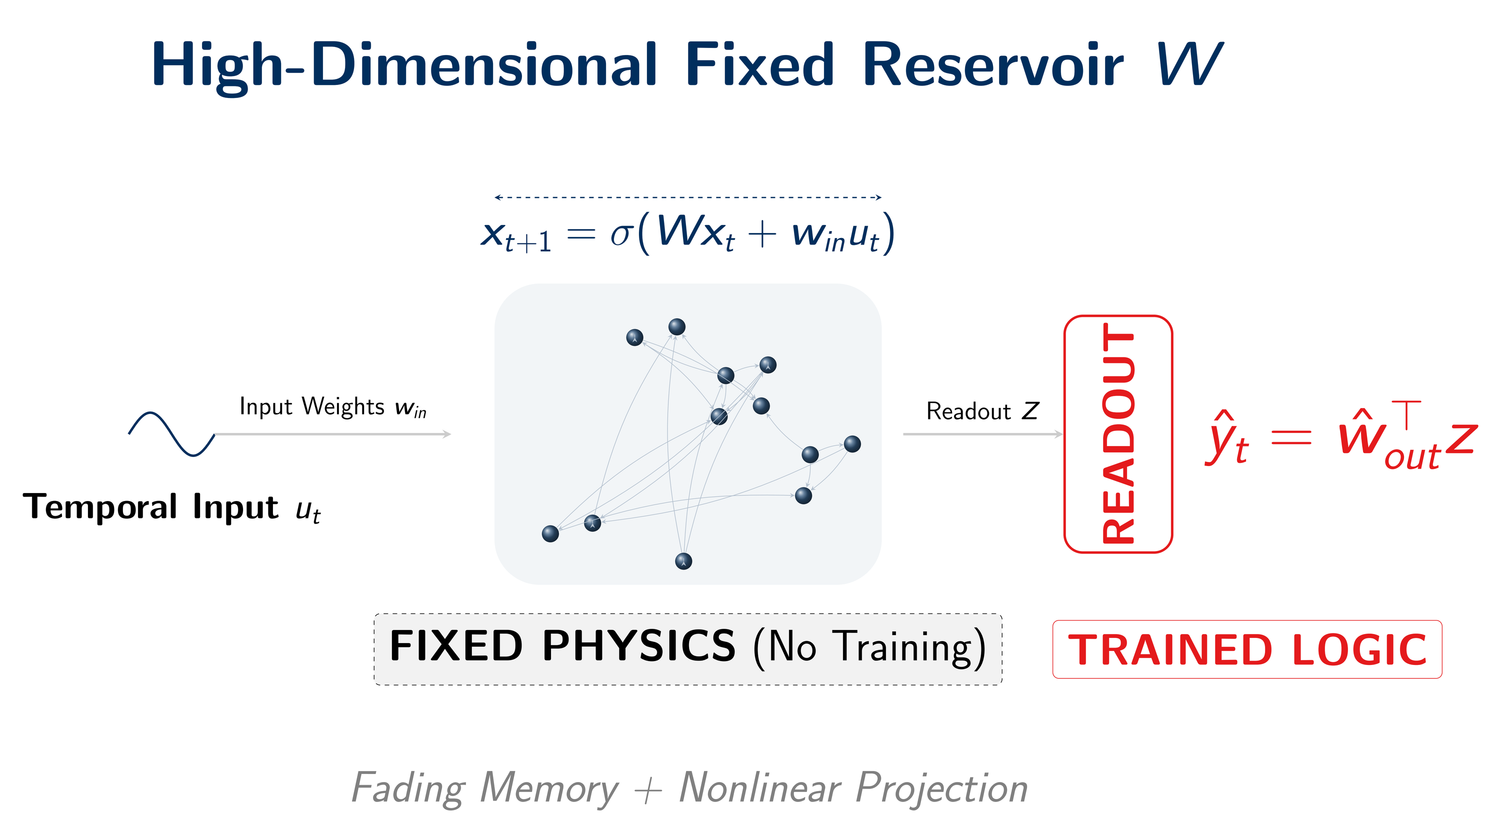

Gap: Despite empirical success, ESNs are still "black boxes."

Our Contribution: Use **Random Matrix Theory (RMT)** to provide the first precise, interpretable characterization of ESN generalization.

2. The Teacher-Student Framework

Learning task: competition between a latent **Teacher** (Oracle) and a **Student** model trained on N independent samples:

1. **Teacher**: For each sample $i \in \{1,\dots,N\}$, teacher generates a target y_i from an input signal $\mathbf{u}_i \in \mathbb{R}^T$:

$$y_i = \boldsymbol{\theta}_*^\top \mathbf{u}_i + \epsilon_i, \quad \epsilon_i \sim \mathcal{N}(0,\sigma^2) \quad \text{(Measurement Noise)}$$

2. **Student Feature Extraction**: maps each $\mathbf{u}_i$ to a feature $\mathbf{z}_i = F(\mathbf{u}_i) \in \mathbb{R}^n$
   - in the case of **Ridge Regression**: Identity map, $\mathbf{z}_i = \mathbf{u}_i$ (and $n = T$).
   - in the case of **Non-Linear ESN**: Features are the final recurrence state $\mathbf{z}_i = \mathbf{x}_{i,T}$:

$$\mathbf{x}_{i,t+1} = \sigma(\mathbf{W}\mathbf{x}_{i,t} + \mathbf{w}_{in}u_{i,t}) \quad (\mathbf{W}, \mathbf{w}_{in} \text{ are fixed random weights})$$

   with $\sigma(\cdot)$ the **activation function** (e.g., tanh).

3. **Readout Training**: Let $\mathbf{Z} = [\mathbf{z}_1,\dots,\mathbf{z}_N] \in \mathbb{R}^{n\times N}$ be the collection of features and $\mathbf{y} = [y_1,\dots,y_N]^\top \in \mathbb{R}^N$ be the labels. Student learns the readout $\hat{\mathbf{w}}_{out}$:

$$\hat{\mathbf{w}}_{out} = \left(\frac{1}{N}\mathbf{Z}\mathbf{Z}^\top + \lambda\mathbf{I}_n\right)^{-1}\frac{1}{N}\mathbf{Z}\mathbf{y} \in \mathbb{R}^n$$

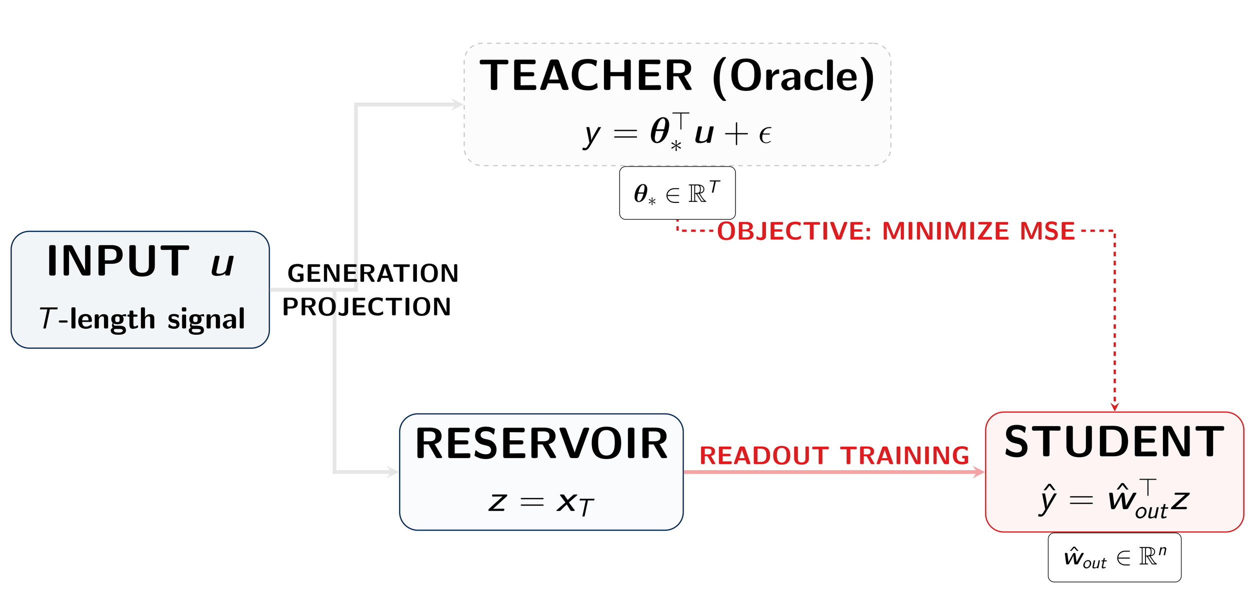

3. The RMT Framework

We characterize performance via the **Out-of-Sample Risk** $\mathcal{R}$, the expected MSE on unseen test data $(\mathbf{u}', y')$.

1. **High-Dimensional Asymptotics**:

$$n, N \to \infty, \quad \frac{n}{N} \to \gamma \in (0,\infty) \quad \text{(Complexity Ratio)}$$

2. **Feature Map**: (q-Exponential) Concentration Reservoir features $\mathbf{z}$ concentrate around their mean through Lipschitz transformations (Gaussian, spherical, complex & close-to-real image, e.g., GAN-generated data).

Deterministic Equivalents allows to "replace" random training matrices with deterministic ones defined using scalars (δ, α):

$$\delta = \frac{1}{N}\mathrm{Tr}(\boldsymbol{\Sigma}_z\bar{\mathbf{Q}}), \quad \alpha = \frac{1}{N}\left\|\frac{\boldsymbol{\Sigma}_z}{1+\delta}\bar{\mathbf{Q}}\right\|_F^2, \quad \bar{\mathbf{Q}} = \left(\frac{\boldsymbol{\Sigma}_z}{1+\delta} + \lambda\mathbf{I}_n\right)^{-1}$$

$\boldsymbol{\Sigma}_z = \mathbb{E}[\mathbf{z}\mathbf{z}^\top]$: Reservoir Covariance (that encodes recurrence).

4. Bias–Variance & Optimal Tuning

Theorem 1 (Exact Asymptotic Risk)

$$\mathcal{R} = \mathcal{B}^2 + \mathcal{V} + \sigma^2$$

via the eigenpairs $\{(\mu_t, \mathbf{v}_t)\}_{t=1}^T$ of $\boldsymbol{\Sigma}_u^{1/2}\mathrm{diag}(\varphi^{t-T})\boldsymbol{\Sigma}_u^{1/2}$:

- **Bias** $\mathcal{B}^2$: Alignment between Teacher signal $\boldsymbol{\theta}_*$ and Reservoir modes $\mathbf{v}_t$.

$$\mathcal{B}^2 = \frac{(1+\delta)^2}{1-\alpha}\sum_{t=1}^T \frac{\lambda^2}{(\mu_t + \lambda(1+\delta))^2}(\boldsymbol{\theta}_*^\top\boldsymbol{\Sigma}_u^{1/2}\mathbf{v}_t)^2$$

  Intuition: Bias is minimized when the reservoir's "echo" $\mathbf{v}_t$ matches the teacher's signal $\boldsymbol{\theta}_*$.
- **Variance** $\mathcal{V} = \frac{\sigma^2\alpha}{1-\alpha}$: Sensitivity to noise σ^2 and the complexity ratio γ.

Implication (Optimal Regularization) For isotropic inputs ($\boldsymbol{\Sigma}_u = \mathbf{I}$), optimal penalty given by

$$\lambda_* = \frac{T}{N}\cdot\mathrm{SNR}, \quad \mathrm{SNR} = \frac{\|\boldsymbol{\theta}_*\|^2}{\sigma^2}$$

independent of reservoir weights, with leaking rate φ handling memory depth.

5. Why ESNs Avoid "Double Descent"

Standard Ridge Regression suffers from **Double Descent**: generalization error spike at $T \approx N$.

- **Double Descent**: As $T/N \to 1$, feature matrix becomes ill-conditioned, leading to variance explosion ($\alpha \to 1$).
- **ESN Advantage**: leaking rate $\varphi < 1$ prevents double descent through implicit bias: approximately low-rank structure with $\alpha < 1$.

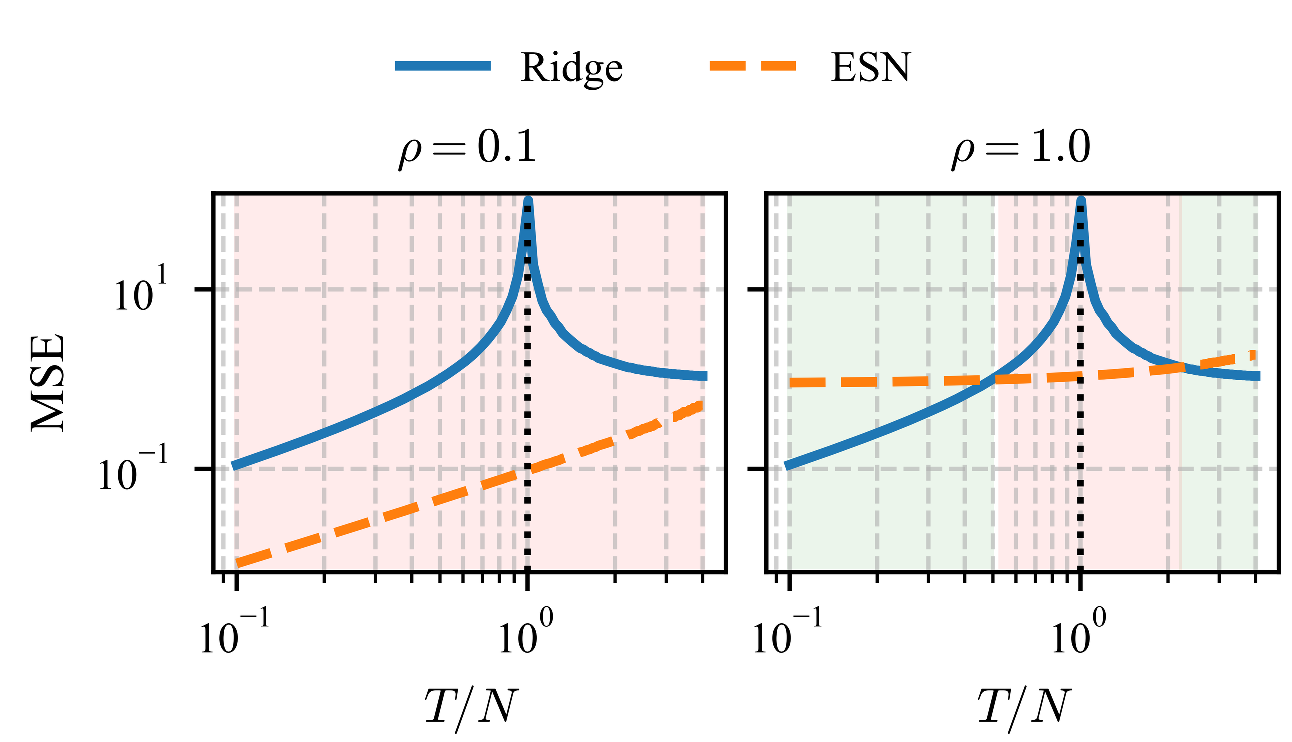

Left: Short memory ($\rho \ll 1$); Right: Long memory ($\rho \approx 1$). ESN (red) stays stable while Ridge (blue) spikes.

6. Experiments: Performance & Tuning

1. **When do ESNs beat Ridge Regression?** ESNs work better in **data-scarce** ($N \ll T$) and **short-memory** regimes.
   - **Inductive Bias**: ESNs assume relevant features are in the recent past.
   - **Ridge Failure**: Without temporal structure, Ridge overfits to noise in old inputs when N is small.

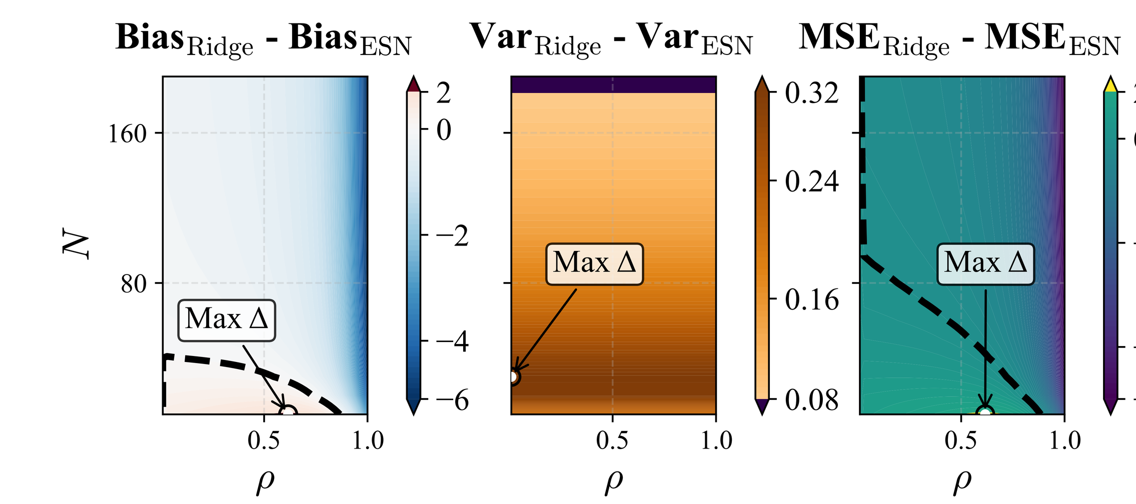

The dashed line shows the boundary where ESN outperforms Ridge.

2. **Validation of the Optimal Regularization Law**: testing the theoretical $\lambda_* = \frac{T}{N}\cdot\mathrm{SNR}$ against empirical grid searches.
   - **Result**: theoretical prediction (dashed vertical line) aligns perfectly with the minimum of the empirical error curve.
   - **Impact**: zero-shot hyperparameter tuning now possible for Linear ESNs.

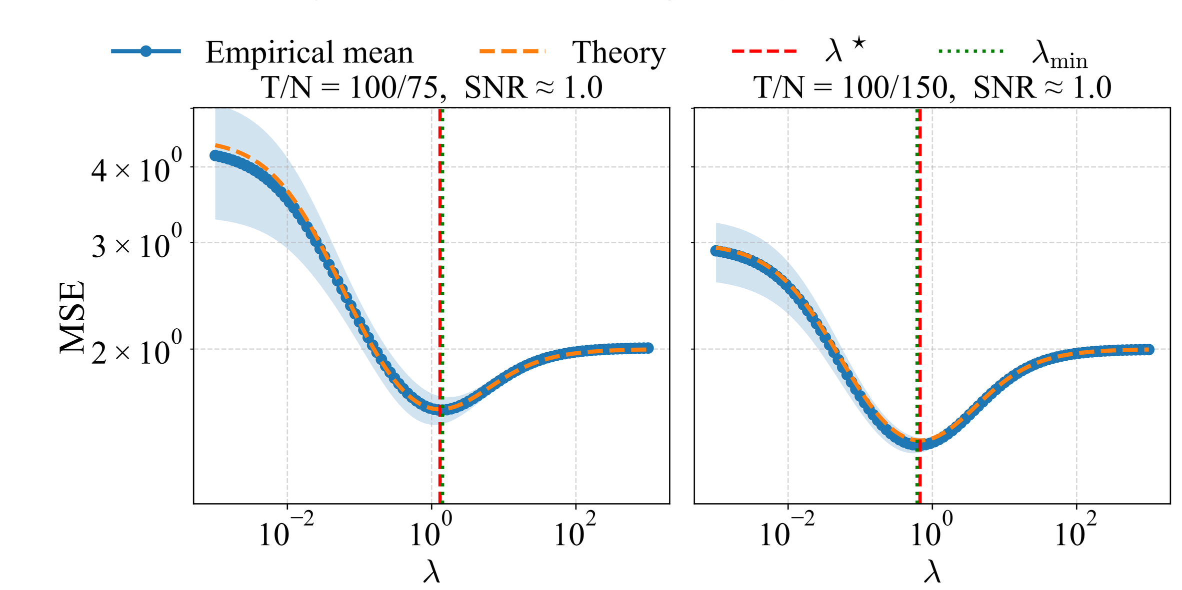

7. Conclusion & Future Outlook

Summary of Contributions:

- First **precise bias-variance formulas** for ESNs using RMT.
- **Absence of double descent** as a consequence of eigenspectral properties.
- Provide a **universal tuning law** λ_* for optimal generalization.

Future Work: Extending this framework to **nonlinear activations** $\sigma(\cdot)$ and the dynamics of deeper RNNs/LSTMs/SSMs.

Code & Paper: `github.com/maliktiomoko/ESN_theory`